

Capitalized Title Here

by Author One, Author Two, Author Three

Abstract An abstract of less than 150 words.

TODO:

- Order of authorship?
- Review Q-Q plots and P-P plots, including other arrangements, and what is implemented in other packages
- Write about what the package implements
- Give examples
 - Heike: BRFSS example
- Intro/conclusion
- Abstract

Introduction

Univariate distributional assessment is a common thread throughout statistical analyses during both the exploratory and confirmatory stages. When we begin exploring a new data set we often consider the distribution of individual variables before moving on to explore multivariate relationships. After a model has been fit to a data set, we must assess whether the distributional assumptions made were reasonable, and if they are not we then must understand the impact this has on the conclusions. Graphics provide arguably the most common way to carry out these univariate assessments. While there are many graphical methods that can be used for distribution exploration and assessment, probability plotting is one of the most common graphical approaches used.

Probability plotting refers to a family of methods based on the cumulative distribution function (CDF), most notably quantile (Q-Q) plots and probability (P-P) plots (Wilk and Gnanadesikan, 1968). In this paper, we focus on comparing an empirical distribution to a theoretical distribution. Let Y_1, \dots, Y_n denote a random sample from an unknown population, and let $\hat{F}_y(q)$ be the empirical cumulative distribution obtained from the sample. Further, let $F(q)$ denote the CDF of a proposed distribution for the sample. A Q-Q plot is constructed by plotting the quantiles of the empirical distribution, $q_y(p) = F_y^{-1}(p)$, against the corresponding quantiles of the theoretical distribution, $q(p) = F^{-1}(p)$. This construction is illustrated in Figure 1. A P-P plot is constructed by plotting $F(q)$ against $\hat{F}_y(q)$ for various quantiles, q . This construction is illustrated in Figure 3. Regardless of the plot constructed, if the two distributions are identical, then the scattplots will be linear with slope 1 and intercept 0. Additionally, Q-Q plots are invariant to linear transformations, so if two random variables differ by a linear transformation a Q-Q plot showing draws from their distributions will still be linear, but with a different slope and intercept, as seen in Figure 1. P-P plots are sensitive to linear transformations.




Figure 1: Illustrating what quantities are being plotted for Q-Q plots.




Figure 2: Illustrating what quantities are being plotted for P-P plots.




Figure 3: Illustrating different designs of probability plots.

While the basic form of both the Q-Q and P-P plots is a scatterplot, additional graphical elements are often added to aid in distributional assessment. For Q-Q plots, a reference line is often drawn through the points $(q(.25), q_y(.25))$ and $(q(.75), q_y(.75))$. For P-P plots a reference line with slope 1 and intercept 0 is used. In both plots, pointwise or simultaneous confidence bands are often added around the reference line to further aid in the visual assessment.

Innovations to Q-Q and P-P plots have also been proposed. [Loy et al. \(2016\)](#) discuss the creation of detrended Q-Q plots, where the y -axis is changed to show the difference between q_y and the reference line. Consequently, the line representing the agreement with the theoretical distribution is the x -axis. [Loy et al. \(2016\)](#) find that detrended Q-Q plots are more powerful than other designs, so long as the y -axis limits are set so that the aspect ratio is kept the same as in the traditional Q-Q plot. In reliability and survival analysis, probability plots often refer to a hybrid probability plot, where the CDF of the proposed theoretical distribution is plotted against the empirical order statistics, and transformations are applied to each axis to linearize the CDF (cf. [Meeker and Escobar, 1998](#), chapter 6). This hybrid probability plot is invariant to linear transformations.

Q-Q plots have been implemented in various forms in R, but none provide a complete implementation of the probability plotting framework. Normal quantile plots, where a sample is compared to the standard normal distribution, are implemented using the `qqplot` and `qqline` in **base** graphics ([R Core Team, 2012](#)). `qqmath` in **lattice** provides a general framework for Q-Q plots, comparing a sample to any theoretical distribution by specifying the quantile function ([Sarkar, 2008](#)). `qqPlot` in the **car** package also allows for the assessment of non-normal distribution and adds pointwise confidence bands based on the standard errors of the order statistics or the parametric bootstrap ([Fox and Weisberg, 2011](#)). **ggplot2** provides `geom_qq` and `geom_qq_line`, enabling the creation of traditional Q-Q plots with a reference line, much like those created using `qqmath`. **qqplotr** extends `\pkg{ggplot2}` to provide the most complete implementation of probability plotting.

In the remainder of this paper, we introduce the probability plotting framework provided by **qqplotr**. . . FILL THIS IN ONCE OTHER SECTIONS ARE WRITTEN. . .

TODO: FIGURE OUT WHERE TO INTRODUCE TS BANDS ([Aldor-Noiman et al., 2013](#))

Implementing probability plots in the ggplot2 framework

The qqplotr package extends some ggplot2 functionalities by permitting the drawing of both quantile-quantile (Q-Q) and probability-probability (P-P) points, lines, and confidence bands. The functions of this package also allow a detrend adjustment of the plots, proposed by Thode (2002) to help reduce visual bias when assessing the results.

The functions of this package, implemented as Stats from ggplot2, are divided into two groups: (1) Q-Q and (2) P-P plot.

Both groups are composed of three functions: *point*, *line*, and *band*. Those Stats complement each other when drawn together, but they may also be plotted independently.

Below we will give an overview of all those Stats and, further in the document, we will present some usage examples.

Q-Q plots

- `stat_qq_point` This is a modified version of `ggplot2::stat_qq` with some parameters adjustments and a new option to detrend the points.
- `stat_qq_line` Draws a reference line based on the data quantiles, as in `stats::qqline`.
- `stat_qq_band` Draws confidence bands based on three methods: "normal", "bs", and "ts":
 - "normal" constructs simultaneous confidence bands based on Normal confidence intervals;
 - "bs" creates pointwise confidence bands based on a parametric bootstrap;
 - "ts" constructs tail-sensitive confidence bands, as proposed by Aldor-Noiman et al. (2013).

In order to facilitate the visualization of multiple Q-Q band methods at the same time, the `geom_qq_band` Geom was also implemented. Its usage will be illustrated further below.

P-P plots

- `stat_pp_point` Plots cumulative probabilities versus probability points. The cumulative probability function is constructed with the sample data, and then evaluated at each probability point.
- `stat_pp_line` Draws a reference identity line ($x = y$).
- `stat_pp_band` Draws confidence bands. For now, only the bootstrap version ("bs") is available.

Examples

In this section, we demonstrate the capabilities of the **qqplotr** package.

```
library(qqplotr)
```

BRFSS example

The Center for Disease Control and Prevention runs an annual telephone survey, the Behavioral Risk Factor Surveillance System (BRFSS), to keep track of the US populations' 'health-related risk behaviors, chronic health conditions, and use of preventive services'.

Close to half a million interviews are conducted each year. Here, we are focussing on the 2012 responses for Iowa. 7166 responses were gathered across 359 questions and derived variables. Among these, are people's height and weight, which we are going to assess in more detail.

Figure 4 shows two Q-Q plots side by side. For each of the plots, a sample of 200 men and 200 women is drawn from the overall number of responses. On the left hand side, individuals' heights are plotted in a Q-Q plot comparing raw heights to a normal distribution. We see that the distributions for both men and women (colour) is showing horizontal steps: this indicates that the distributional assessment is heavily dominated by the discreteness in the data, as most survey participants responded to the question of their height to the closest inch. On the right hand side of figure 4, we use jittering; this means that we add a random number generated from a random uniform distribution on ± 0.5 inch to the reported height. By this mean we extinguish any effect that discreteness might have on the distribution. This brings the observed distribution much closer to a normal distribution. Note that separate normal distributions were fitted for each gender, not surprisingly, the resulting distributions have different means (women are on average 6 inch shorter than men in this dataset). Interestingly, the slope of the two genders is similar, indicating that the same

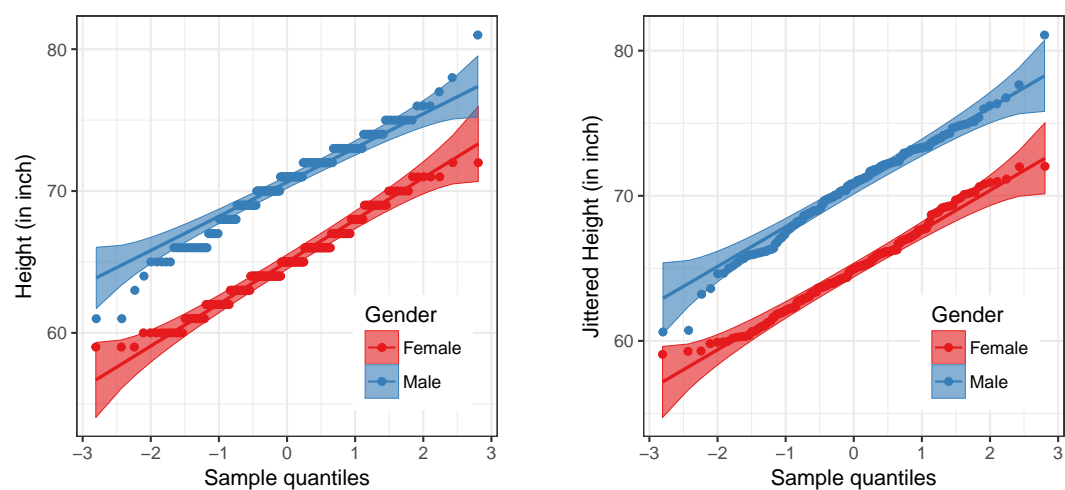


Figure 4: Sample (200 men and 200 women) of raw heights (left) and jittered heights (right). The distribution on the left is dominated by the discreteness of the data. On the right we see that except for some outliers an assumption of normality for people’s height is not completely absurd.

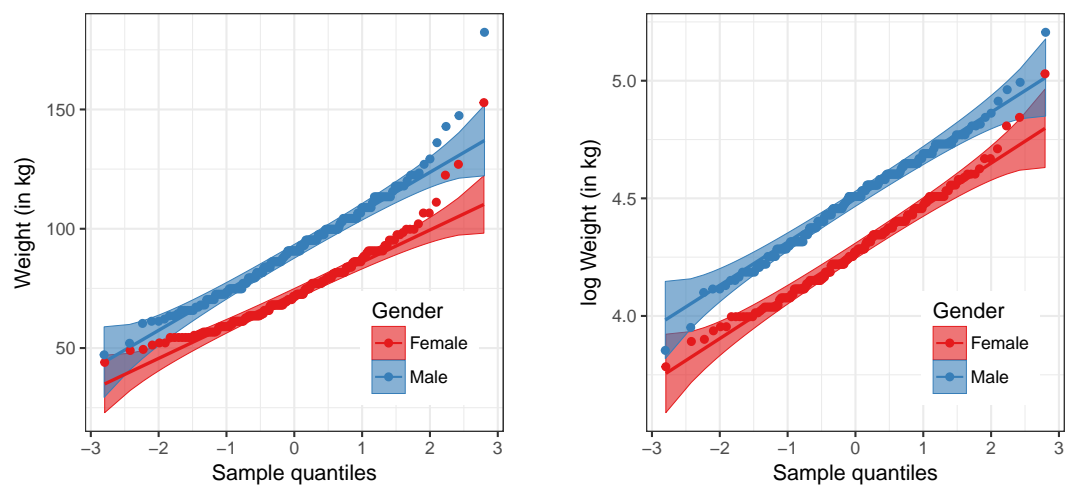


Figure 5: Sample (200 men and 200 women) of weights. Unlike people’s height, weight seems to be heavily right skewed with some additional outliers on the extreme left. (left plot), on the right, weight was log-transformed before its distribution is compared to a theoretical normal.

scale parameter fits both genders’ distributions (the standard deviation of height in the data set is 2.97 inch for men and 2.91 inch for women).

Unlike respondents’ heights, their weights do not seem to be normally distributed. Figure 5 shows again two Q-Q plots. The Q-Q plot on the left uses raw weights and compares to a normal distribution. From the curved points we see that tails of the observed distribution are heavier than expected under a normal distribution. On the right, weights are log-transformed. We see that a normal distribution for each of the genders shows –with the exceptions of a few extreme outliers– a reasonable fit.

Instead of transforming the observed values, we can change the theoretical distribution against which we compare. Figure 6 shows two Q-Q plots where a log-normal distribution is chosen as the theoretical distribution. On the left, we compare against a log-normal distribution with mean 4.389 and standard deviation 0.223 (the log-transformed averages of average weight and standard deviation in Iowa’s population). Again, the fits seem reasonable. On the right, parameters for the log-normal distribution are fit separately. The fits are slightly different from

Table 1: this table is just for us at the moment

SEX	mean_wt	mean_log_wt	sd_wt	sd_log_wt
1	91.50342	4.495766	19.22981	0.2011216
2	74.37085	4.282780	17.90126	0.2254663

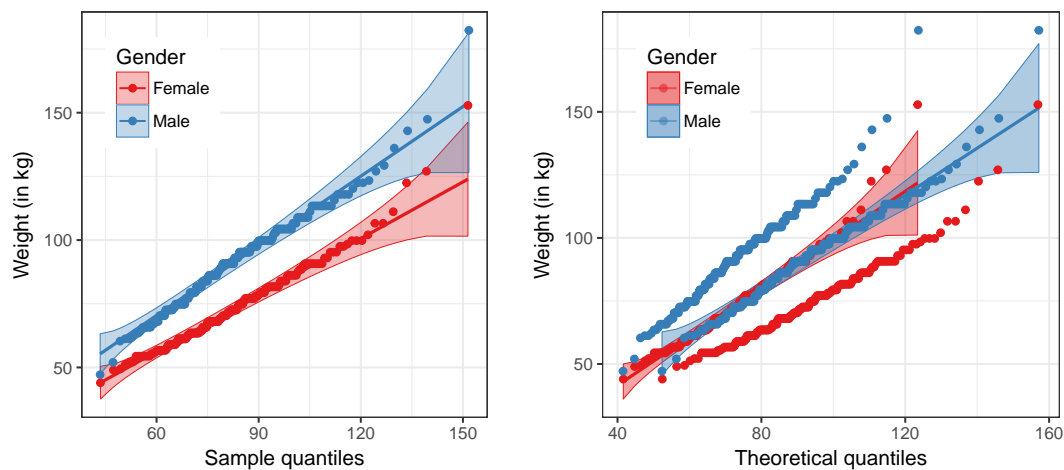


Figure 6: Sample (200 men and 200 women) of weights. On the left, the theoretical distribution is changed to a log normal. On the right, we additionally estimate shift and scale parameters for each of the genders separately before comparing distributions to a log-normal. XXX need to fix alpha

Using a non-standard distribution

ADAM: PULL DATA OFF EXTERNAL HARD DRIVE, EXPLAIN THE EXAMPLE...

The “standard” distributions provided in `pkg{stats}` do not include a left-skewed distribution, which is clearly needed to model the long jump heights. To handle this type of situation, **qqplotr** let’s users create probability plots for all of the distributions defined in the **stats** package as well as other sources, so long as the distributions are defined following the conventions laid out in the **stats** package. Specifically, for some distribution there must be density/mass (d prefix), CDF (p prefix), quantile (q prefix), and simulation (r prefix) functions.

To illustrate this process we will consider the smallest extreme value (SEV) distribution. This is a left-skewed distribution, which is a key feature needed in order to model the long jump results. Below, we define the suite of distribution functions needed to utilize the SEV distribution.

```
# CDF
psev <- function(q, mu, sigma) {
  z <- (x - mu) / sigma
  1 - exp(-exp(z))
}

# PDF
dsev <- function(x, mu, sigma) {
  z <- (x - mu) / sigma
  (1 / sigma) * exp(z - exp(z))
}

# Quantile function
qsev <- function(p, mu, sigma) {
  mu + log(-log(1 - p)) * sigma
}

# Simulation function
rsev <- function(n, mu, sigma) {
  qsev(runif(n))
}
```

Summary

Write this section once the rest of the paper is done.

Bibliography

- S. Aldor-Noiman, L. D. Brown, A. Buja, W. Rolke, and R. A. Stine. The power to see: A new graphical test of normality. *The American Statistician*, 67(4):249–260, 1 Nov. 2013. [p2]
- J. Fox and S. Weisberg. *An R Companion to Applied Regression*. Sage, Thousand Oaks CA, second edition, 2011. URL <http://socserv.socsci.mcmaster.ca/jfox/Books/Companion>. [p2]
- A. Loy, L. Follett, and H. Hofmann. Variations of Q–Q plots: The power of our eyes! *The American Statistician*, 70(2):202–214, 2 Apr. 2016. [p2]
- W. Q. Meeker and L. A. Escobar. *Statistical methods for reliability data*. John Wiley & Sons, 1998. [p2]
- R Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria, 2012. URL <http://www.R-project.org/>. ISBN 3-900051-07-0. [p2]
- D. Sarkar. *Lattice: Multivariate Data Visualization with R*. Springer, New York, 2008. URL <http://lmdvr.r-forge.r-project.org>. ISBN 978-0-387-75968-5. [p2]
- M. B. Wilk and R. Gnanadesikan. Probability plotting methods for the analysis of data. *Biometrika*, 55(1):1–17, Mar. 1968. [p1]

Acknowledgements

Mention GSoC here...

Author One

Affiliation

line 1

line 2

[author1@work](#)

Author Two

Affiliation

line 1

line 2

[author2@work](#)

Author Three

Affiliation

line 1

line 2

[author3@work](#)